

```

import torch
x = torch.arange(12, dtype=torch.float32)
x
tensor([ 0.,  1.,  2.,  3.,  4.,  5.,  6.,  7.,  8.,  9., 10., 11.])
x.numel()
12
x.shape
torch.Size([12])
X = x.reshape(3,4)
X
tensor([[ 0.,  1.,  2.,  3.],
        [ 4.,  5.,  6.,  7.],
        [ 8.,  9., 10., 11.]])
X = x.reshape(-1,4)
X
tensor([[ 0.,  1.,  2.,  3.],
        [ 4.,  5.,  6.,  7.],
        [ 8.,  9., 10., 11.]])
X = x.reshape(3,-1)
X
tensor([[ 0.,  1.,  2.,  3.],
        [ 4.,  5.,  6.,  7.],
        [ 8.,  9., 10., 11.]])
torch.zeros((2,3,4))
tensor([[[0., 0., 0., 0.],
        [0., 0., 0., 0.],
        [0., 0., 0., 0.]],

        [[0., 0., 0., 0.],
        [0., 0., 0., 0.],
        [0., 0., 0., 0.]])
torch.ones((2,3,4))
tensor([[[1., 1., 1., 1.],
        [1., 1., 1., 1.],
        [1., 1., 1., 1.]],

        [[1., 1., 1., 1.],

```

```

        [1., 1., 1., 1.],
        [1., 1., 1., 1.]])
torch.randn(3,4)
tensor([[ 0.2522,  0.3935, -0.4037,  1.4716],
        [-1.4807, -0.4549,  0.4695,  0.8155],
        [ 0.9276, -1.4810, -0.2608,  0.4720]])
torch.tensor([[2,1,4,3],[1,2,3,4],[4,3,2,1]])
tensor([[2, 1, 4, 3],
        [1, 2, 3, 4],
        [4, 3, 2, 1]])
X
tensor([[ 0.,  1.,  2.,  3.],
        [ 4.,  5.,  6.,  7.],
        [ 8.,  9., 10., 11.]])
X[-1], X[1:3]
(tensor([ 8.,  9., 10., 11.]),
 tensor([[ 4.,  5.,  6.,  7.],
        [ 8.,  9., 10., 11.])))
X[1, 2] = 17
X
tensor([[ 0.,  1.,  2.,  3.],
        [ 4.,  5., 17.,  7.],
        [ 8.,  9., 10., 11.]])
X[:,2,:] = 12
X
tensor([[12., 12., 12., 12.],
        [12., 12., 12., 12.],
        [ 8.,  9., 10., 11.]])
x
tensor([ 0.,  1.,  2.,  3.,  4.,  5.,  6.,  7.,  8.,  9., 10., 11.])
torch.exp(x)
tensor([1.0000e+00, 2.7183e+00, 7.3891e+00, 2.0086e+01, 5.4598e+01,
        1.4841e+02,
        4.0343e+02, 1.0966e+03, 2.9810e+03, 8.1031e+03, 2.2026e+04,
        5.9874e+04])

```

```

x = torch.tensor([1.0,2,4,8])
y = torch.tensor([2,2,2,2])
x+y, x-y, x*y, x/y, x**y

(tensor([ 3.,  4.,  6., 10.]),
 tensor([-1.,  0.,  2.,  6.]),
 tensor([ 2.,  4.,  8., 16.]),
 tensor([0.5000, 1.0000, 2.0000, 4.0000]),
 tensor([ 1.,  4., 16., 64.]))

X = torch.arange(12, dtype=torch.float32).reshape((3,4))
Y = torch.tensor([[2.0,1,4,3],[1,2,3,4],[4,3,2,1]])
torch.cat((X,Y),dim=0), torch.cat((X,Y),dim=1)

(tensor([[ 0.,  1.,  2.,  3.],
         [ 4.,  5.,  6.,  7.],
         [ 8.,  9., 10., 11.],
         [ 2.,  1.,  4.,  3.],
         [ 1.,  2.,  3.,  4.],
         [ 4.,  3.,  2.,  1.]]),
 tensor([[ 0.,  1.,  2.,  3.,  2.,  1.,  4.,  3.],
         [ 4.,  5.,  6.,  7.,  1.,  2.,  3.,  4.],
         [ 8.,  9., 10., 11.,  4.,  3.,  2.,  1.])))

X, Y

(tensor([[ 0.,  1.,  2.,  3.],
         [ 4.,  5.,  6.,  7.],
         [ 8.,  9., 10., 11.]]),
 tensor([[2., 1., 4., 3.],
         [1., 2., 3., 4.],
         [4., 3., 2., 1.])))

X == Y

tensor([[False,  True, False,  True],
        [False, False, False, False],
        [False, False, False, False]])

X.sum()

tensor(66.)

a = torch.arange(3).reshape((3,1))
b = torch.arange(2).reshape((1,2))
a,b

(tensor([[0],
         [1],
         [2]]),
 tensor([[0, 1]]))

```

a+b

```
tensor([[0, 1],  
        [1, 2],  
        [2, 3]])
```

```
before = id(Y)  
Y = Y + X  
id(Y) == before
```

False

```
Z = torch.zeros_like(Y)  
print('id(Z):', id(Z))  
Z[:] = X + Y  
print('id(Z):', id(Z))
```

```
id(Z): 2332161146952  
id(Z): 2332161146952
```

```
before = id(X)  
X += Y  
id(X) == before
```

True

```
A = X.numpy()  
A
```

```
array([[ 2.,  3.,  8.,  9.],  
       [ 9., 12., 15., 18.],  
       [20., 21., 22., 23.]], dtype=float32)
```

```
B = torch.from_numpy(A)  
B
```

```
tensor([[ 2.,  3.,  8.,  9.],  
        [ 9., 12., 15., 18.],  
        [20., 21., 22., 23.]])
```

```
type(A), type(B)
```

```
(numpy.ndarray, torch.Tensor)
```

```
a = torch.tensor([3.5])  
a, a.item(), float(a), int(a)
```

```
(tensor([3.5000]), 3.5, 3.5, 3)
```

```
X = torch.arange(12, dtype=torch.float32).reshape((3,4))  
Y = torch.tensor([[2.0, 1, 4, 3], [1, 2, 3, 4], [4, 3, 2, 1]])  
X, Y
```

```
(tensor([[ 0.,  1.,  2.,  3.],
         [ 4.,  5.,  6.,  7.],
         [ 8.,  9., 10., 11.]]),
 tensor([[2., 1., 4., 3.],
         [1., 2., 3., 4.],
         [4., 3., 2., 1.])))
```

$X == Y$, $X < Y$, $X > Y$

```
(tensor([[False,  True, False,  True],
         [False, False, False, False],
         [False, False, False, False]]),
 tensor([[ True, False,  True, False],
         [False, False, False, False],
         [False, False, False, False]]),
 tensor([[False, False, False, False],
         [ True,  True,  True,  True],
         [ True,  True,  True,  True]]))
```

```
a = torch.arange(6).reshape((2,3,1))
a
```

```
tensor([[[0],
         [1],
         [2]],
        [[3],
         [4],
         [5]]])
```

```
b = torch.arange(2).reshape((1,2))
b
```

```
tensor([[0, 1]])
```

a+b

```
tensor([[[0, 1],
         [1, 2],
         [2, 3]],
        [[3, 4],
         [4, 5],
         [5, 6]]])
```

```
import os

os.makedirs(os.path.join('.', 'data'), exist_ok=True)
data_file = os.path.join('.', 'data', 'house_tiny.csv')
with open(data_file, 'w') as f:
    f.write(''NumRooms,RoofType,Price
NA,NA,127500
2,NA,106000
4,Slate,178100
NA,NA,140000'')
```

```
import pandas as pd
```

```
data = pd.read_csv(data_file)
print(data)
```

	NumRooms	RoofType	Price
0	NaN	NaN	127500
1	2.0	NaN	106000
2	4.0	Slate	178100
3	NaN	NaN	140000

```
inputs = data.iloc[:, 0:2]
inputs = pd.get_dummies(inputs, dummy_na=True)
print(inputs)
```

	NumRooms	RoofType_Slate	RoofType_nan
0	NaN	0	1
1	2.0	0	1
2	4.0	1	0
3	NaN	0	1

```
inputs = inputs.fillna(inputs.mean())
print(inputs)
```

	NumRooms	RoofType_Slate	RoofType_nan
0	3.0	0	1
1	2.0	0	1
2	4.0	1	0
3	3.0	0	1

```
targets = data.iloc[:, 2]
targets
```

0	127500
1	106000
2	178100
3	140000

Name: Price, dtype: int64

```
import torch
```

```

X = torch.tensor(inputs.to_numpy(dtype=float))
y = torch.tensor(targets.to_numpy(dtype=float))
X,y

(tensor([[3., 0., 1.],
        [2., 0., 1.],
        [4., 1., 0.],
        [3., 0., 1.]], dtype=torch.float64),
 tensor([127500., 106000., 178100., 140000.], dtype=torch.float64))

Y = y.reshape((4,1))
Y

tensor([[127500.],
        [106000.],
        [178100.],
        [140000.]], dtype=torch.float64)

data = pd.read_csv('wdbc.data')
print(data)

```

```

      842302  M  17.99  10.38  122.8   1001  0.1184  0.2776
0.3001  \
0      842517  M  20.57  17.77  132.90  1326.0  0.08474  0.07864
0.08690
1      84300903  M  19.69  21.25  130.00  1203.0  0.10960  0.15990
0.19740
2      84348301  M  11.42  20.38   77.58   386.1  0.14250  0.28390
0.24140
3      84358402  M  20.29  14.34  135.10  1297.0  0.10030  0.13280
0.19800
4      843786  M  12.45  15.70   82.57   477.1  0.12780  0.17000
0.15780
...      ... ..      ...      ...      ...      ...      ...
...
563      926424  M  21.56  22.39  142.00  1479.0  0.11100  0.11590
0.24390
564      926682  M  20.13  28.25  131.20  1261.0  0.09780  0.10340
0.14400
565      926954  M  16.60  28.08  108.30   858.1  0.08455  0.10230
0.09251
566      927241  M  20.60  29.33  140.10  1265.0  0.11780  0.27700
0.35140
567      92751  B   7.76  24.54   47.92   181.0  0.05263  0.04362
0.00000

      0.1471  ...   25.38  17.33  184.6   2019  0.1622  0.6656
0.7119  \
0      0.07017  ...   24.990  23.41  158.80  1956.0  0.12380  0.18660
0.2416

```

```

1    0.12790 ... 23.570 25.53 152.50 1709.0 0.14440 0.42450
0.4504
2    0.10520 ... 14.910 26.50 98.87 567.7 0.20980 0.86630
0.6869
3    0.10430 ... 22.540 16.67 152.20 1575.0 0.13740 0.20500
0.4000
4    0.08089 ... 15.470 23.75 103.40 741.6 0.17910 0.52490
0.5355
...
...
563 0.13890 ... 25.450 26.40 166.10 2027.0 0.14100 0.21130
0.4107
564 0.09791 ... 23.690 38.25 155.00 1731.0 0.11660 0.19220
0.3215
565 0.05302 ... 18.980 34.12 126.70 1124.0 0.11390 0.30940
0.3403
566 0.15200 ... 25.740 39.42 184.60 1821.0 0.16500 0.86810
0.9387
567 0.00000 ... 9.456 30.37 59.16 268.6 0.08996 0.06444
0.0000

```

```

    0.2654 0.4601 0.1189
0    0.1860 0.2750 0.08902
1    0.2430 0.3613 0.08758
2    0.2575 0.6638 0.17300
3    0.1625 0.2364 0.07678
4    0.1741 0.3985 0.12440
..
563 0.2216 0.2060 0.07115
564 0.1628 0.2572 0.06637
565 0.1418 0.2218 0.07820
566 0.2650 0.4087 0.12400
567 0.0000 0.2871 0.07039

```

[568 rows x 32 columns]

```
Thresh=max(data.isnull().sum(axis=0))
```

```
print(Thresh)
```

```
pro_data=data.dropna(axis=1,thresh=data.shape[0]-Thresh+1)
```

```
print(pro_data)
```

```
0
```

```
Empty DataFrame
```

```
Columns: []
```

```
Index: [0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17,
18, 19, 20, 21, 22, 23, 24, 25, 26, 27, 28, 29, 30, 31, 32, 33, 34,
35, 36, 37, 38, 39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51,
52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64, 65, 66, 67, 68,
69, 70, 71, 72, 73, 74, 75, 76, 77, 78, 79, 80, 81, 82, 83, 84, 85,
86, 87, 88, 89, 90, 91, 92, 93, 94, 95, 96, 97, 98, 99, ...]
```



```
[568 rows x 0 columns]
```

```
targets = data.iloc[:,2]  
targets
```

```
      842302  M  
0      842517  M  
1  84300903  M  
2  84348301  M
```

```

import torch
x = torch.tensor(3.0)
y = torch.tensor(2.0)
x, y, x+y, x*y, x/y, x**y

(tensor(3.), tensor(2.), tensor(5.), tensor(6.), tensor(1.5000),
tensor(9.))

x = torch.arange(3)
x

tensor([0, 1, 2])

x[2]

tensor(2)

len(x)

3

x.shape

torch.Size([3])

A = torch.arange(6).reshape(3, 2)
A

tensor([[0, 1],
        [2, 3],
        [4, 5]])

A.T

tensor([[0, 2, 4],
        [1, 3, 5]])

A = torch.tensor([[1, 2, 3], [2, 0, 4], [3, 4, 5]])
A == A.T

tensor([[True, True, True],
        [True, True, True],
        [True, True, True]])

len(A), A.shape

(3, torch.Size([3, 3]))

A = torch.tensor([[1, 2, 3], [2, 0, 4], [3, 4, 5], [8,8,8]])
len(A)

4

torch.arange(24).reshape(2,3,4)

```

```
tensor([[[ 0,  1,  2,  3],
          [ 4,  5,  6,  7],
          [ 8,  9, 10, 11]],

        [[12, 13, 14, 15],
          [16, 17, 18, 19],
          [20, 21, 22, 23]]])
```

```
A = torch.arange(6, dtype=torch.float32).reshape(2,3)
B = A.clone()
A,B,A+B
```

```
(tensor([[0., 1., 2.],
          [3., 4., 5.]]),
 tensor([[0., 1., 2.],
          [3., 4., 5.]]),
 tensor([[ 0.,  2.,  4.],
          [ 6.,  8., 10.]])
```

*A*B # elementwise product*

```
tensor([[ 0.,  1.,  4.],
        [ 9., 16., 25.]])
```

```
a = 2
X = torch.arange(24).reshape(2,3,4)
X, a*X
```

```
(tensor([[[ 0,  1,  2,  3],
          [ 4,  5,  6,  7],
          [ 8,  9, 10, 11]],

        [[12, 13, 14, 15],
          [16, 17, 18, 19],
          [20, 21, 22, 23]]]),
 tensor([[[ 2,  3,  4,  5],
          [ 6,  7,  8,  9],
          [10, 11, 12, 13]],

        [[14, 15, 16, 17],
          [18, 19, 20, 21],
          [22, 23, 24, 25]]]))
```

```
(a*X), (a*X).shape
```

```
(tensor([[[ 0,  2,  4,  6],
          [ 8, 10, 12, 14],
          [16, 18, 20, 22]],

        [[24, 26, 28, 30],
          [32, 34, 36, 38]]])
```

```

        [40, 42, 44, 46]]]),
    torch.Size([2, 3, 4]))

x = torch.arange(3, dtype=torch.float32)
x, x.sum()

(tensor([0., 1., 2.]), tensor(3.))

A, A.shape, A.sum(axis=0), A.sum(axis=0).shape #axis 0 is row 0

(tensor([[0., 1., 2.],
        [3., 4., 5.]]),
    torch.Size([2, 3]),
    tensor([3., 5., 7.]),
    torch.Size([3]))

A, A.shape, A.sum(axis=1), A.sum(axis=1).shape #axis 0 is row 0

(tensor([[0., 1., 2.],
        [3., 4., 5.]]),
    torch.Size([2, 3]),
    tensor([ 3., 12.]),
    torch.Size([2]))

A.sum(axis=[0, 1]), A.sum(), A.sum(axis=[0, 1]) == A.sum()

(tensor(15.), tensor(15.), tensor(True))

A.mean(), A.sum(), A.numel(), A.sum()/A.numel() # numel = num of
elements

(tensor(2.5000), tensor(15.), 6, tensor(2.5000))

A, A.mean(axis=0), A.sum(axis=0), A.shape[0], A.sum(axis=0) /
A.shape[0]

(tensor([[0., 1., 2.],
        [3., 4., 5.]]),
    tensor([1.5000, 2.5000, 3.5000]),
    tensor([3., 5., 7.]),
    2,
    tensor([1.5000, 2.5000, 3.5000]))

sum_A = A.sum(axis=1, keepdims=True)
A, A.shape, sum_A, sum_A.shape

(tensor([[0., 1., 2.],
        [3., 4., 5.]]),
    torch.Size([2, 3]),
    tensor([[ 3.],
        [12.]]),
    torch.Size([2, 1]))

```

```

A.sum(axis=1), A.sum(axis=1).shape
(tensor([ 3., 12.]), torch.Size([2]))
A, sum_A, A / sum_A
(tensor([[0., 1., 2.],
         [3., 4., 5.]]),
 tensor([[ 3.],
         [12.]]),
 tensor([[0.0000, 0.3333, 0.6667],
         [0.2500, 0.3333, 0.4167]]))

A.cumsum(axis=0)
(tensor([[0., 1., 2.],
         [3., 5., 7.]])

y = torch.ones(3, dtype = torch.float32)
x, y, torch.dot(x, y)
(tensor([0., 1., 2.]), tensor([1., 1., 1.]), tensor(3.))
x*y, torch.sum(x * y)
(tensor([0., 1., 2.]), tensor(3.))

A.shape, x.shape, torch.mv(A, x), A@x
(torch.Size([2, 3]), torch.Size([3]), tensor([ 5., 14.]), tensor([ 5.,
14.]))

A, x, A@x
(tensor([[0., 1., 2.],
         [3., 4., 5.]]),
 tensor([0., 1., 2.]),
 tensor([ 5., 14.]))

B = torch.ones(3, 4)
A, B, torch.mm(A, B), A@B
(tensor([[0., 1., 2.],
         [3., 4., 5.]]),
 tensor([[1., 1., 1., 1.],
         [1., 1., 1., 1.],
         [1., 1., 1., 1.]]),
 tensor([[ 3.,  3.,  3.,  3.],
         [12., 12., 12., 12.]]),
 tensor([[ 3.,  3.,  3.,  3.],
         [12., 12., 12., 12.]])

```

```
u = torch.tensor([3.0, -4.0])
torch.norm(u)

tensor(5.)

torch.abs(u).sum()

tensor(7.)

torch.ones((4, 9)), torch.norm(torch.ones((4, 9)))

(tensor([[1., 1., 1., 1., 1., 1., 1., 1., 1.],
         [1., 1., 1., 1., 1., 1., 1., 1., 1.],
         [1., 1., 1., 1., 1., 1., 1., 1., 1.],
         [1., 1., 1., 1., 1., 1., 1., 1., 1.]]),
 tensor(6.))
```

```

import torch
x = torch.arange(4.0)
x
tensor([0., 1., 2., 3.])
x.requires_grad_(True)
x.grad
y = 2 * torch.dot(x,x)
y
tensor(28., grad_fn=<MulBackward0>)
y.backward()
x.grad
tensor([ 0.,  4.,  8., 12.])
x.grad == 4 * x
tensor([True, True, True, True])
x.grad.zero_() # Reset the gradient
y = x.sum()
y.backward()
x.grad
tensor([1., 1., 1., 1.])
x.grad.zero_()
y = x * x
y.backward(gradient=torch.ones(len(y))) # Faster: y.sum().backward()
x.grad
tensor([0., 2., 4., 6.])
x.grad.zero_()
y = x * x
u = y.detach()
z = u * x
z.sum().backward()
x.grad == u
tensor([True, True, True, True])
x.grad.zero_()
y.sum().backward()
x.grad == 2 * x
tensor([True, True, True, True])

```

```
def f(a):  
    b = a * 2  
    while b.norm() < 1000:  
        b = b * 2  
    if b.sum() > 0:  
        c = b  
    else:  
        c = 100 * b  
    return c  
  
a = torch.randn(size=(), requires_grad=True)  
d = f(a)  
d.backward()  
  
a.grad == d / a
```



```

%matplotlib inline
import math
import time
import numpy as np
import torch
from d2l import torch as d2l

n = 10000
a = torch.ones(n)
b = torch.ones(n)
c = torch.zeros(n)
t = time.time()
for i in range(n):
    c[i] = a[i] + b[i]
f'{time.time() - t:.5f} sec'

'0.09100 sec'

t = time.time()
d = a + b
f'{time.time() - t:.5f} sec'

'0.00000 sec'

def normal(x, mu, sigma):
    p = 1 / math.sqrt(2 * math.pi * sigma**2)
    return p * np.exp(-0.5 * (x - mu)**2 / sigma**2)

# Use NumPy again for visualization
x = np.arange(-7, 7, 0.01)

# Mean and standard deviation pairs
params = [(0, 1), (0, 2), (3, 1)]
d2l.plot(x, [normal(x, mu, sigma) for mu, sigma in params],
xlabel='x',
ylabel='p(x)', figsize=(4.5, 2.5),
legend=[f'mean {mu}, std {sigma}' for mu, sigma in params])

```

```
# Use NumPy again for visualization
x = np.arange(-7, 7, 0.01)

# Mean and standard deviation pairs
params = [(0, 2), (1, 3), (3, 1)]
d2l.plot(x, [normal(x, mu, sigma) for mu, sigma in params],
xlabel='x',
        ylabel='p(x)', figsize=(4.5, 2.5),
        legend=[f'mean {mu}, std {sigma}' for mu, sigma in params])
```

```

import time
import numpy as np
import torch
from torch import nn
from d2l import torch as d2l

def add_to_class(Class): #@save
    """Register functions as methods in created class."""
    def wrapper(obj):
        setattr(Class, obj.__name__, obj)
    return wrapper

class A:
    def __init__(self):
        self.b = 1

a = A()

@add_to_class(A)
def do(self):
    print('Class attribute "b" is', self.b)

a.do()

Class attribute "b" is 1

class HyperParameters: #@save
    """The base class of hyperparameters."""
    def save_hyperparameters(self, ignore=[]):
        raise NotImplemented

import time
import numpy as np
import torch
from torch import nn
from d2l import torch as d2l

class B(d2l.HyperParameters):
    def __init__(self, a, b, c):
        self.save_hyperparameters(ignore=['c'])
        print('self.a =', self.a, 'self.b =', self.b)
        print('There is no self.c =', not hasattr(self, 'c'))

b = B(a=1, b=2, c=3)

self.a = 1 self.b = 2
There is no self.c = True

class ProgressBoard(d2l.HyperParameters): #@save
    """The board that plots data points in animation."""
    def __init__(self, xlabel=None, ylabel=None, xlim=None,

```

```

        ylim=None, xscale='linear', yscale='linear',
        ls=['-', '--', '-.', ':'], colors=['C0', 'C1', 'C2',
'C3'],
        fig=None, axes=None, figsize=(3.5, 2.5),
display=True):
    self.save_hyperparameters()

    def draw(self, x, y, label, every_n=1):
        raise NotImplemented

board = d2l.ProgressBoard('x')
for x in np.arange(0, 10, 0.1):
    board.draw(x, np.sin(x), 'sin', every_n=2)
    board.draw(x, np.cos(x), 'cos', every_n=10)

```

```

board = d2l.ProgressBoard('x')
for x in np.arange(0, 10, 0.1):
    board.draw(x, np.sin(2*x), 'sin', every_n=2)
    board.draw(x, np.cos(x/2), 'cos', every_n=2)

```

```

class Module(nn.Module, d2l.HyperParameters):  #@save
    """The base class of models."""
    def __init__(self, plot_train_per_epoch=2,
plot_valid_per_epoch=1):
        super().__init__()
        self.save_hyperparameters()
        self.board = ProgressBoard()

    def loss(self, y_hat, y):
        raise NotImplementedError

    def forward(self, X):
        assert hasattr(self, 'net'), 'Neural network is defined'
        return self.net(X)

    def plot(self, key, value, train):
        """Plot a point in animation."""
        assert hasattr(self, 'trainer'), 'Trainer is not initied'
        self.board.xlabel = 'epoch'
        if train:
            x = self.trainer.train_batch_idx / \
                self.trainer.num_train_batches
            n = self.trainer.num_train_batches / \
                self.plot_train_per_epoch
        else:
            x = self.trainer.epoch + 1
            n = self.trainer.num_val_batches / \
                self.plot_valid_per_epoch
        self.board.draw(x, value.to(d2l.cpu()).detach().numpy(),
                        ('train_' if train else 'val_') + key,
                        every_n=int(n))

    def training_step(self, batch):
        l = self.loss(self(*batch[:-1]), batch[-1])
        self.plot('loss', l, train=True)
        return l

    def validation_step(self, batch):
        l = self.loss(self(*batch[:-1]), batch[-1])
        self.plot('loss', l, train=False)

    def configure_optimizers(self):
        raise NotImplementedError

class DataModule(d2l.HyperParameters):  #@save
    """The base class of data."""
    def __init__(self, root='../data', num_workers=4):
        self.save_hyperparameters()

    def get_dataloader(self, train):

```

```

        raise NotImplementedError

    def train_dataloader(self):
        return self.get_dataloader(train=True)

    def val_dataloader(self):
        return self.get_dataloader(train=False)

class Trainer(d2l.HyperParameters):  #@save
    """The base class for training models with data."""
    def __init__(self, max_epochs, num_gpus=0, gradient_clip_val=0):
        self.save_hyperparameters()
        assert num_gpus == 0, 'No GPU support yet'

    def prepare_data(self, data):
        self.train_dataloader = data.train_dataloader()
        self.val_dataloader = data.val_dataloader()
        self.num_train_batches = len(self.train_dataloader)
        self.num_val_batches = (len(self.val_dataloader)
                                if self.val_dataloader is not None
                                else 0)

    def prepare_model(self, model):
        model.trainer = self
        model.board.xlim = [0, self.max_epochs]
        self.model = model

    def fit(self, model, data):
        self.prepare_data(data)
        self.prepare_model(model)
        self.optim = model.configure_optimizers()
        self.epoch = 0
        self.train_batch_idx = 0
        self.val_batch_idx = 0
        for self.epoch in range(self.max_epochs):
            self.fit_epoch()

    def fit_epoch(self):
        raise NotImplementedError

```

```

%matplotlib inline
import torch
from d2l import torch as d2l

class LinearRegressionScratch(d2l.Module): #@save
    """The linear regression model implemented from scratch."""
    def __init__(self, num_inputs, lr, sigma=0.01):
        super().__init__()
        self.save_hyperparameters()
        self.w = torch.normal(0, sigma, (num_inputs, 1),
requires_grad=True)
        self.b = torch.zeros(1, requires_grad=True)

@d2l.add_to_class(LinearRegressionScratch) #@save
def forward(self, X):
    return torch.matmul(X, self.w) + self.b

@d2l.add_to_class(LinearRegressionScratch) #@save
def loss(self, y_hat, y):
    l = (y_hat - y) ** 2 / 2
    return l.mean()

class SGD(d2l.HyperParameters): #@save
    """Minibatch stochastic gradient descent."""
    def __init__(self, params, lr):
        self.save_hyperparameters()

    def step(self):
        for param in self.params:
            param -= self.lr * param.grad

    def zero_grad(self):
        for param in self.params:
            if param.grad is not None:
                param.grad.zero_()

@d2l.add_to_class(LinearRegressionScratch) #@save
def configure_optimizers(self):
    return SGD([self.w, self.b], self.lr)

@d2l.add_to_class(d2l.Trainer) #@save
def prepare_batch(self, batch):
    return batch

@d2l.add_to_class(d2l.Trainer) #@save
def fit_epoch(self):
    self.model.train()
    for batch in self.train_dataloader:
        loss = self.model.training_step(self.prepare_batch(batch))
        self.optim.zero_grad()
        with torch.no_grad():

```

```

        loss.backward()
        if self.gradient_clip_val > 0: # To be discussed later
            self.clip_gradients(self.gradient_clip_val,
self.model)
            self.optim.step()
            self.train_batch_idx += 1
        if self.val_dataloader is None:
            return
        self.model.eval()
        for batch in self.val_dataloader:
            with torch.no_grad():
                self.model.validation_step(self.prepare_batch(batch))
            self.val_batch_idx += 1

model = LinearRegressionScratch(2, lr=0.03)
data = d2l.SyntheticRegressionData(w=torch.tensor([2, -3.4]), b=4.2)
trainer = d2l.Trainer(max_epochs=3)
trainer.fit(model, data)

```

```

with torch.no_grad():
    print(f'error in estimating w: {data.w -
model.w.reshape(data.w.shape)}')
    print(f'error in estimating b: {data.b - model.b}')

error in estimating w: tensor([ 0.1082, -0.2121])
error in estimating b: tensor([0.2522])

model = LinearRegressionScratch(2, lr=0.03)
data = d2l.SyntheticRegressionData(w=torch.tensor([2, -3.4]), b=4.2)
trainer = d2l.Trainer(max_epochs=5)
trainer.fit(model, data)

```



```
with torch.no_grad():  
    print(f'error in estimating w: {data.w -  
model.w.reshape(data.w.shape)}')  
    print(f'error in estimating b: {data.b - model.b}')  
  
error in estimating w: tensor([ 0.0116, -0.0380])  
error in estimating b: tensor([0.0293])
```

```

%matplotlib inline
import time
import torch
import torchvision
from torchvision import transforms
from d2l import torch as d2l

d2l.use_svg_display()

class FashionMNIST(d2l.DataModule):
    def __init__(self, batch_size=64, resize=(28,28)):
        super().__init__()
        self.save_hyperparameters()
        trans = transforms.Compose([transforms.Resize(resize),
                                     transforms.ToTensor()])
        self.train = torchvision.datasets.FashionMNIST(
            root=self.root, train=True, transform=trans,
            download=True)
        self.val = torchvision.datasets.FashionMNIST(
            root=self.root, train=False, transform=trans,
            download=True)

data = FashionMNIST(resize=(32, 32))
len(data.train), len(data.val)

(60000, 10000)

data.train[0][0].shape
torch.Size([1, 32, 32])

@d2l.add_to_class(FashionMNIST)  #@save
def text_labels(self, indices):
    """Return text labels."""
    labels = ['t-shirt', 'trouser', 'pullover', 'dress', 'coat',
              'sandal', 'shirt', 'sneaker', 'bag', 'ankle boot']
    return [labels[int(i)] for i in indices]

@d2l.add_to_class(FashionMNIST)
def get_dataloader(self, train):
    data = self.train if train else self.val
    return torch.utils.data.DataLoader(data, self.batch_size,
    shuffle=train, num_workers=self.num_workers)

X, y = next(iter(data.train_dataloader()))
print(X.shape, X.dtype, y.shape, y.dtype)

torch.Size([64, 1, 32, 32]) torch.float32 torch.Size([64]) torch.int64

tic = time.time()
for X, y in data.train_dataloader():

```

```
        continue
    f'{time.time() - tic:.2f} sec'

'3.79 sec'

def show_images(imgs, num_rows, num_cols, titles=None, scale=1.5):
    #@save
    """Plot a list of images."""
    raise NotImplementedError

@d2l.add_to_class(FashionMNIST)
def visualize(self, batch, nrows=1, ncols=8, labels=[]):
    x, y = batch
    if not labels:
        labels = self.text_labels(y)
    d2l.show_images(x.squeeze(1), nrows, ncols, titles=labels)

batch = next(iter(data.val_dataloader()))
data.visualize(batch)
```

```

import torch
from d2l import torch as d2l

class Classifier(d2l.Module):
    def validation_step(self, batch):
        Y_hat = self(*batch[:-1])
        self.plot('loss', self.loss(Y_hat, batch[-1]), train=False)
        self.plot('acc', self.accuracy(Y_hat, batch[-1]), train=False)

@d2l.add_to_class(d2l.Module)
def configure_optimizers(self):
    return torch.optim.SGD(self.parameters(), lr=self.lr)

@d2l.add_to_class(Classifier)
def accuracy(self, Y_hat, Y, averaged=True):
    """Compute the number of correct predictions."""
    Y_hat = Y_hat.reshape((-1, Y_hat.shape[-1]))
    preds = Y_hat.argmax(axis=1).type(Y.dtype)
    compare = (preds == Y.reshape(-1)).type(torch.float32)
    return compare.mean() if averaged else compare

```

```

import torch
from d2l import torch as d2l
X = torch.tensor([[1.0, 2.0, 3.0], [4.0, 5.0, 6.0]])
X.sum(0, keepdims=True), X.sum(1, keepdims=True)

(tensor([[5., 7., 9.]]),
 tensor([[ 6.],
         [15.])))

def softmax(X):
    X_exp = torch.exp(X)
    partition = X_exp.sum(1, keepdims=True)
    return X_exp / partition

X = torch.rand((2, 5))
X_prob = softmax(X)
X_prob, X_prob.sum(1)

(tensor([[0.1918, 0.2102, 0.1433, 0.2283, 0.2264],
         [0.2267, 0.2162, 0.1267, 0.2004, 0.2300]]),
 tensor([1., 1.]))

class SoftmaxRegressionScratch(d2l.Classifier):
    def __init__(self, num_inputs, num_outputs, lr, sigma=0.01):
        super().__init__()
        self.save_hyperparameters()
        self.W = torch.normal(0, sigma, size=(num_inputs,
num_outputs),
                                requires_grad=True)
        self.b = torch.zeros(num_outputs, requires_grad=True)

    def parameters(self):
        return [self.W, self.b]

@d2l.add_to_class(SoftmaxRegressionScratch)
def forward(self, X):
    X = X.reshape((-1, self.W.shape[0]))
    return softmax(torch.matmul(X, self.W) + self.b)

y = torch.tensor([0, 2])
y_hat = torch.tensor([[0.1, 0.3, 0.6], [0.3, 0.2, 0.5]])
y_hat[[0, 1], y]

tensor([0.1000, 0.5000])

def cross_entropy(y_hat, y):
    return -torch.log(y_hat[list(range(len(y_hat))), y]).mean()

cross_entropy(y_hat, y)

tensor(1.4979)

```

```

@d2l.add_to_class(SoftmaxRegressionScratch)
def loss(self, y_hat, y):
    return cross_entropy(y_hat, y)

data = d2l.FashionMNIST(batch_size=256)
model = SoftmaxRegressionScratch(num_inputs=784, num_outputs=10,
    lr=0.1)
trainer = d2l.Trainer(max_epochs=10)
trainer.fit(model, data)

```

```

X, y = next(iter(data.val_dataloader()))
preds = model(X).argmax(axis=1)
preds.shape

torch.Size([256])

wrong = preds.type(y.dtype) != y
X, y, preds = X[wrong], y[wrong], preds[wrong]
labels = [a+'\n'+b for a, b in zip(
    data.text_labels(y), data.text_labels(preds))]
data.visualize([X, y], labels=labels)

```

```

import torch
from torch import nn
from d2l import torch as d2l

class MLPScratch(d2l.Classifier):
    def __init__(self, num_inputs, num_outputs, num_hiddens, lr,
sigma=0.01):
        super().__init__()
        self.save_hyperparameters()
        self.W1 = nn.Parameter(torch.randn(num_inputs, num_hiddens) *
sigma)
        self.b1 = nn.Parameter(torch.zeros(num_hiddens))
        self.W2 = nn.Parameter(torch.randn(num_hiddens, num_outputs) *
sigma)
        self.b2 = nn.Parameter(torch.zeros(num_outputs))

    def relu(X):
        a = torch.zeros_like(X)
        return torch.max(X, a)

@d2l.add_to_class(MLPScratch)
def forward(self, X):
    X = X.reshape((-1, self.num_inputs))
    H = relu(torch.matmul(X, self.W1) + self.b1)
    return torch.matmul(H, self.W2) + self.b2

model = MLPScratch(num_inputs=784, num_outputs=10, num_hiddens=256,
lr=0.1)
data = d2l.FashionMNIST(batch_size=256)
trainer = d2l.Trainer(max_epochs=10)
trainer.fit(model, data)

```

```

class MLP(d2l.Classifier):
    def __init__(self, num_outputs, num_hiddens, lr):
        super().__init__()

```

```
        self.save_hyperparameters()
        self.net = nn.Sequential(nn.Flatten(),
nn.LazyLinear(num_hiddens),
                                nn.ReLU(),
nn.LazyLinear(num_outputs))
model = MLP(num_outputs=10, num_hiddens=256, lr=0.1)
trainer.fit(model, data)
```