

```
In [3]: import torch
        from torch import nn
        from d2l import torch as d2l
```

```
In [4]: def corr2d(X, K): #@save
        """Compute 2D cross-correlation."""
        h, w = K.shape
        Y = torch.zeros((X.shape[0] - h + 1, X.shape[1] - w + 1))
        for i in range(Y.shape[0]):
            for j in range(Y.shape[1]):
                Y[i, j] = (X[i:i + h, j:j + w] * K).sum()
        return Y
```

```
In [5]: X = torch.tensor([[0.0, 1.0, 2.0], [3.0, 4.0, 5.0], [6.0, 7.0, 8.0]])
        K = torch.tensor([[0.0, 1.0], [2.0, 3.0]])
        corr2d(X, K)
```

```
Out[5]: tensor([[19., 25.],
                [37., 43.]])
```

```
In [6]: class Conv2D(nn.Module):
        def __init__(self, kernel_size):
            super().__init__()
            self.weight = nn.Parameter(torch.rand(kernel_size))
            self.bias = nn.Parameter(torch.zeros(1))

        def forward(self, x):
            return corr2d(x, self.weight) + self.bias
```

```
In [7]: X = torch.ones((6, 8))
        X[:, 2:6] = 0
        X
```

```
Out[7]: tensor([[1., 1., 0., 0., 0., 0., 1., 1.],
                [1., 1., 0., 0., 0., 0., 1., 1.],
                [1., 1., 0., 0., 0., 0., 1., 1.],
                [1., 1., 0., 0., 0., 0., 1., 1.],
                [1., 1., 0., 0., 0., 0., 1., 1.],
                [1., 1., 0., 0., 0., 0., 1., 1.]])
```

```
In [8]: K = torch.tensor([[1.0, -1.0]])
        Y = corr2d(X, K)
        Y
```

```
Out[8]: tensor([[ 0.,  1.,  0.,  0.,  0., -1.,  0.],
                [ 0.,  1.,  0.,  0.,  0., -1.,  0.],
                [ 0.,  1.,  0.,  0.,  0., -1.,  0.],
                [ 0.,  1.,  0.,  0.,  0., -1.,  0.],
                [ 0.,  1.,  0.,  0.,  0., -1.,  0.],
                [ 0.,  1.,  0.,  0.,  0., -1.,  0.]])
```

```
In [9]: corr2d(X.t(), K)
```

```
Out[9]: tensor([[0., 0., 0., 0., 0.],
                [0., 0., 0., 0., 0.],
                [0., 0., 0., 0., 0.],
                [0., 0., 0., 0., 0.],
                [0., 0., 0., 0., 0.],
                [0., 0., 0., 0., 0.],
                [0., 0., 0., 0., 0.],
                [0., 0., 0., 0., 0.]])
```

```
In [10]: # Construct a two-dimensional convolutional layer with 1 output channel and  
# kernel of shape (1, 2). For the sake of simplicity, we ignore the bias here  
conv2d = nn.LazyConv2d(1, kernel_size=(1, 2), bias=False)
```

```
# The two-dimensional convolutional layer uses four-dimensional input and  
# output in the format of (example, channel, height, width), where the batch  
# size (number of examples in the batch) and the number of channels are both  
X = X.reshape((1, 1, 6, 8))  
Y = Y.reshape((1, 1, 6, 7))  
lr = 3e-2 # Learning rate
```

```
for i in range(10):  
    Y_hat = conv2d(X)  
    l = (Y_hat - Y) ** 2  
    conv2d.zero_grad()  
    l.sum().backward()  
    # Update the kernel  
    conv2d.weight.data[:] -= lr * conv2d.weight.grad  
    if (i + 1) % 2 == 0:  
        print(f'epoch {i + 1}, loss {l.sum():.3f}')
```

```
epoch 2, loss 12.772  
epoch 4, loss 2.285  
epoch 6, loss 0.441  
epoch 8, loss 0.098  
epoch 10, loss 0.026
```

```
C:\Users\sjkim\AppData\Local\Programs\Python\Python310\lib\site-packages\torch\nn\modules\lazy.py:180: UserWarning: Lazy modules are a new feature under heavy development so changes to the API or functionality can happen at any moment.
```

```
warnings.warn('Lazy modules are a new feature under heavy development ')
```

```
In [11]: conv2d.weight.data.reshape((1, 2))
```

```
Out[11]: tensor([[ 0.9951, -0.9686]])
```

```
In [ ]:
```

```
In [1]: import torch
        from torch import nn
```

```
In [4]: def comp_conv2d(conv2d, X):
        X = X.reshape((1, 1) + X.shape)
        Y = conv2d(X)
        # Strip the first two dimensions: examples and channels
        return Y.reshape(Y.shape[2:])

conv2d = nn.LazyConv2d(1, kernel_size=3, padding=1)
X = torch.rand(size=(8, 8))
comp_conv2d(conv2d, X).shape
```

```
Out[4]: torch.Size([8, 8])
```

```
In [5]: conv2d = nn.LazyConv2d(1, kernel_size=(5, 3), padding=(2, 1))
        comp_conv2d(conv2d, X).shape
```

```
Out[5]: torch.Size([8, 8])
```

```
In [6]: conv2d = nn.LazyConv2d(1, kernel_size=3, padding=1, stride=2)
        comp_conv2d(conv2d, X).shape
```

```
Out[6]: torch.Size([4, 4])
```

```
In [7]: conv2d = nn.LazyConv2d(1, kernel_size=(3, 5), padding=(0, 1), stride=(3, 4))
        comp_conv2d(conv2d, X).shape
```

```
Out[7]: torch.Size([2, 2])
```

```
In [ ]:
```

```
In [1]: import torch
        from d2l import torch as d2l
```

```
In [2]: def corr2d_multi_in(X, K):
        return sum(d2l.corr2d(x, k) for x, k in zip(X, K))
```

```
In [3]: X = torch.tensor([[[0.0, 1.0, 2.0], [3.0, 4.0, 5.0], [6.0, 7.0, 8.0]],
                          [[1.0, 2.0, 3.0], [4.0, 5.0, 6.0], [7.0, 8.0, 9.0]]])
        K = torch.tensor([[[0.0, 1.0], [2.0, 3.0]], [[1.0, 2.0], [3.0, 4.0]]])

        corr2d_multi_in(X, K)
```

```
Out[3]: tensor([[ 56.,  72.],
                [104., 120.]])
```

```
In [4]: def corr2d_multi_in_out(X, K):
        return torch.stack([corr2d_multi_in(X, k) for k in K], 0)
```

```
In [5]: K = torch.stack((K, K + 1, K + 2), 0)
        K.shape
```

```
Out[5]: torch.Size([3, 2, 2, 2])
```

```
In [6]: corr2d_multi_in_out(X, K)
```

```
Out[6]: tensor([[[ 56.,  72.],
                  [104., 120.]],

                [[ 76., 100.],
                  [148., 172.]],

                [[ 96., 128.],
                  [192., 224.]])
```

```
In [7]: def corr2d_multi_in_out_1x1(X, K):
        c_i, h, w = X.shape
        c_o = K.shape[0]
        X = X.reshape((c_i, h * w))
        K = K.reshape((c_o, c_i))
        Y = torch.matmul(K, X)
        return Y.reshape((c_o, h, w))
```

```
In [8]: X = torch.normal(0, 1, (3, 3, 3))
        K = torch.normal(0, 1, (2, 3, 1, 1))
        Y1 = corr2d_multi_in_out_1x1(X, K)
        Y2 = corr2d_multi_in_out(X, K)
        assert float(torch.abs(Y1 - Y2).sum()) < 1e-6
```

```
In [ ]:
```

```
In [1]: import torch
        from torch import nn
        from d2l import torch as d2l
```

```
In [2]: def pool2d(X, pool_size, mode='max'):
        p_h, p_w = pool_size
        Y = torch.zeros((X.shape[0] - p_h + 1, X.shape[1] - p_w + 1))
        for i in range(Y.shape[0]):
            for j in range(Y.shape[1]):
                if mode == 'max':
                    Y[i, j] = X[i: i + p_h, j: j + p_w].max()
                elif mode == 'avg':
                    Y[i, j] = X[i: i + p_h, j: j + p_w].mean()
        return Y
```

```
In [3]: X = torch.tensor([[0.0, 1.0, 2.0], [3.0, 4.0, 5.0], [6.0, 7.0, 8.0]])
        pool2d(X, (2, 2))
```

```
Out[3]: tensor([[4., 5.],
                [7., 8.]])
```

```
In [4]: pool2d(X, (2, 2), 'avg')
```

```
Out[4]: tensor([[2., 3.],
                [5., 6.]])
```

```
In [5]: X = torch.arange(16, dtype=torch.float32).reshape((1, 1, 4, 4))
        X
```

```
Out[5]: tensor([[[[ 0.,  1.,  2.,  3.],
                    [ 4.,  5.,  6.,  7.],
                    [ 8.,  9., 10., 11.],
                    [12., 13., 14., 15.]]]]])
```

```
In [6]: pool2d = nn.MaxPool2d(3)
        pool2d(X)
```

```
Out[6]: tensor([[[[10.]]]])
```

```
In [7]: pool2d = nn.MaxPool2d(3, padding=1, stride=2)
        pool2d(X)
```

```
Out[7]: tensor([[[[ 5.,  7.],
                    [13., 15.]]]]])
```

```
In [8]: pool2d = nn.MaxPool2d((2, 3), stride=(2, 3), padding=(0, 1))
        pool2d(X)
```

```
Out[8]: tensor([[[[ 5.,  7.],
                    [13., 15.]]]]])
```

```
In [9]: X = torch.cat((X, X + 1), 1)
        X
```

```
Out[9]: tensor([[[[ 0.,  1.,  2.,  3.],
                  [ 4.,  5.,  6.,  7.],
                  [ 8.,  9., 10., 11.],
                  [12., 13., 14., 15.]],

                [[ 1.,  2.,  3.,  4.],
                  [ 5.,  6.,  7.,  8.],
                  [ 9., 10., 11., 12.],
                  [13., 14., 15., 16.]]]]])
```

```
In [10]: pool2d = nn.MaxPool2d(3, padding=1, stride=2)
         pool2d(X)
```

```
Out[10]: tensor([[[[ 5.,  7.],
                  [13., 15.]],

                [[ 6.,  8.],
                  [14., 16.]]]]])
```

```
In [ ]:
```

```
In [1]: import torch
        from torch import nn
        from d2l import torch as d2l
```

```
In [2]: def init_cnn(module):
        """Initialize weights for CNNs."""
        if type(module) == nn.Linear or type(module) == nn.Conv2d:
            nn.init.xavier_uniform_(module.weight)

        class LeNet(d2l.Classifier):
            """The LeNet-5 model."""
            def __init__(self, lr=0.1, num_classes=10):
                super().__init__()
                self.save_hyperparameters()
                self.net = nn.Sequential(
                    nn.LazyConv2d(6, kernel_size=5, padding=2), nn.Sigmoid(),
                    nn.AvgPool2d(kernel_size=2, stride=2),
                    nn.LazyConv2d(16, kernel_size=5), nn.Sigmoid(),
                    nn.AvgPool2d(kernel_size=2, stride=2),
                    nn.Flatten(),
                    nn.LazyLinear(120), nn.Sigmoid(),
                    nn.LazyLinear(84), nn.Sigmoid(),
                    nn.LazyLinear(num_classes))
```

```
In [3]: @d2l.add_to_class(d2l.Classifier)  #@save
        def layer_summary(self, X_shape):
            X = torch.randn(*X_shape)
            for layer in self.net:
                X = layer(X)
                print(layer.__class__.__name__, 'output shape:\t', X.shape)

        model = LeNet()
        model.layer_summary((1, 1, 28, 28))
```

```
Conv2d output shape:      torch.Size([1, 6, 28, 28])
Sigmoid output shape:    torch.Size([1, 6, 28, 28])
AvgPool2d output shape:  torch.Size([1, 6, 14, 14])
Conv2d output shape:    torch.Size([1, 16, 10, 10])
Sigmoid output shape:    torch.Size([1, 16, 10, 10])
AvgPool2d output shape:  torch.Size([1, 16, 5, 5])
Flatten output shape:    torch.Size([1, 400])
Linear output shape:     torch.Size([1, 120])
Sigmoid output shape:    torch.Size([1, 120])
Linear output shape:     torch.Size([1, 84])
Sigmoid output shape:    torch.Size([1, 84])
Linear output shape:     torch.Size([1, 10])
```

C:\Users\sjkim\AppData\Local\Programs\Python\Python310\lib\site-packages\torch\nn\modules\lazy.py:180: UserWarning: Lazy modules are a new feature under heavy development so changes to the API or functionality can happen at any moment.


warnings.warn('Lazy modules are a new feature under heavy development ')

```
In [4]: trainer = d2l.Trainer(max_epochs=10, num_gpus=1)
        data = d2l.FashionMNIST(batch_size=128)
```

```

model = LeNet(lr=0.1)
model.apply_init([next(iter(data.get_dataloader(True)))[0]], init_cnn)
trainer.fit(model, data)

```



```


In [5]: class LeNet(d2l.Classifier):
        """The LeNet-5 model."""
        def __init__(self, lr=0.1, num_classes=10):
            super().__init__()
            self.save_hyperparameters()
            self.net = nn.Sequential(
                nn.LazyConv2d(6, kernel_size=5, padding=2), nn.ReLU(),
                nn.MaxPool2d(kernel_size=2, stride=2),
                nn.LazyConv2d(16, kernel_size=5), nn.ReLU(),
                nn.MaxPool2d(kernel_size=2, stride=2),
                nn.Flatten(),
                nn.LazyLinear(120), nn.ReLU(),
                nn.LazyLinear(84), nn.ReLU(),
                nn.LazyLinear(num_classes))

```

```

In [6]: trainer = d2l.Trainer(max_epochs=10, num_gpus=1)
        data = d2l.FashionMNIST(batch_size=128)
        model = LeNet(lr=0.1)
        model.apply_init([next(iter(data.get_dataloader(True)))[0]], init_cnn)
        trainer.fit(model, data)

```





In [ ]:

This notebook was converted to PDF with [convert.ploomber.io](https://convert.ploomber.io)

```
In [1]: import torch
        from torch import nn
        from d2l import torch as d2l
```

```
In [2]: def vgg_block(num_convs, out_channels):
        layers = []
        for _ in range(num_convs):
            layers.append(nn.LazyConv2d(out_channels, kernel_size=3, padding=1))
            layers.append(nn.ReLU())
        layers.append(nn.MaxPool2d(kernel_size=2, stride=2))
        return nn.Sequential(*layers)
```

```
In [3]: class VGG(d2l.Classifier):
        def __init__(self, arch, lr=0.1, num_classes=10):
            super().__init__()
            self.save_hyperparameters()
            conv_blks = []
            for (num_convs, out_channels) in arch:
                conv_blks.append(vgg_block(num_convs, out_channels))
            self.net = nn.Sequential(
                *conv_blks, nn.Flatten(),
                nn.LazyLinear(4096), nn.ReLU(), nn.Dropout(0.5),
                nn.LazyLinear(4096), nn.ReLU(), nn.Dropout(0.5),
                nn.LazyLinear(num_classes))
            self.net.apply(d2l.init_cnn)
```

```
In [4]: VGG(arch=((1, 64), (1, 128), (2, 256), (2, 512), (2, 512))).layer_summary(
        (1, 1, 224, 224))
```

C:\Users\sjkim\AppData\Local\Programs\Python\Python310\lib\site-packages\torch\nn\modules\lazy.py:180: UserWarning: Lazy modules are a new feature under heavy development so changes to the API or functionality can happen at any moment.

warnings.warn('Lazy modules are a new feature under heavy development ')

```
Sequential output shape:      torch.Size([1, 64, 112, 112])
Sequential output shape:      torch.Size([1, 128, 56, 56])
Sequential output shape:      torch.Size([1, 256, 28, 28])
Sequential output shape:      torch.Size([1, 512, 14, 14])
Sequential output shape:      torch.Size([1, 512, 7, 7])
Flatten output shape:         torch.Size([1, 25088])
Linear output shape:          torch.Size([1, 4096])
ReLU output shape:            torch.Size([1, 4096])
Dropout output shape:         torch.Size([1, 4096])
Linear output shape:          torch.Size([1, 4096])
ReLU output shape:            torch.Size([1, 4096])
Dropout output shape:         torch.Size([1, 4096])
Linear output shape:          torch.Size([1, 10])
```

```
In [ ]: model = VGG(arch=((1, 16), (1, 32), (2, 64), (2, 128), (2, 128)), lr=0.01)
        trainer = d2l.Trainer(max_epochs=10, num_gpus=1)
        data = d2l.FashionMNIST(batch_size=128, resize=(224, 224))
        model.apply_init([next(iter(data.get_dataloader(True)))[0]], d2l.init_cnn)
        trainer.fit(model, data)
```

In [ ]:

This notebook was converted to PDF with [convert.ploomber.io](https://convert.ploomber.io)

```
In [1]: import torch
        from torch import nn
        from torch.nn import functional as F
        from d2l import torch as d2l
```

```
In [2]: class Residual(nn.Module):
        def __init__(self, num_channels, use_1x1conv=False, strides=1):
            super().__init__()
            self.conv1 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1,
                                       stride=strides)
            self.conv2 = nn.LazyConv2d(num_channels, kernel_size=3, padding=1)
            if use_1x1conv:
                self.conv3 = nn.LazyConv2d(num_channels, kernel_size=1,
                                           stride=strides)
            else:
                self.conv3 = None
            self.bn1 = nn.LazyBatchNorm2d()
            self.bn2 = nn.LazyBatchNorm2d()

        def forward(self, X):
            Y = F.relu(self.bn1(self.conv1(X)))
            Y = self.bn2(self.conv2(Y))
            if self.conv3:
                X = self.conv3(X)
            Y += X
            return F.relu(Y)
```

```
In [3]: blk = Residual(3)
        X = torch.randn(4, 3, 6, 6)
        blk(X).shape
```

C:\Users\sjkim\AppData\Local\Programs\Python\Python310\lib\site-packages\torch\nn\modules\lazy.py:180: UserWarning: Lazy modules are a new feature under heavy development so changes to the API or functionality can happen at any moment.

warnings.warn('Lazy modules are a new feature under heavy development ')

```
Out[3]: torch.Size([4, 3, 6, 6])
```

```
In [4]: blk = Residual(6, use_1x1conv=True, strides=2)
        blk(X).shape
```

```
Out[4]: torch.Size([4, 6, 3, 3])
```

```
In [5]: class ResNet(d2l.Classifier):
        def b1(self):
            return nn.Sequential(
                nn.LazyConv2d(64, kernel_size=7, stride=2, padding=3),
                nn.LazyBatchNorm2d(), nn.ReLU(),
                nn.MaxPool2d(kernel_size=3, stride=2, padding=1))
```

```
In [6]: @d2l.add_to_class(ResNet)
        def block(self, num_residuals, num_channels, first_block=False):
            blk = []
```

```

    for i in range(num_residuals):
        if i == 0 and not first_block:
            blk.append(Residual(num_channels, use_1x1conv=True, strides=2))
        else:
            blk.append(Residual(num_channels))
    return nn.Sequential(*blk)

```

```

In [7]: @d2l.add_to_class(ResNet)
def __init__(self, arch, lr=0.1, num_classes=10):
    super(ResNet, self).__init__()
    self.save_hyperparameters()
    self.net = nn.Sequential(self.b1())
    for i, b in enumerate(arch):
        self.net.add_module(f'b{i+2}', self.block(*b, first_block=(i==0)))
    self.net.add_module('last', nn.Sequential(
        nn.AdaptiveAvgPool2d((1, 1)), nn.Flatten(),
        nn.Linear(num_classes)))
    self.net.apply(d2l.init_cnn)

```

```

In [8]: class ResNet18(ResNet):
def __init__(self, lr=0.1, num_classes=10):
    super().__init__(((2, 64), (2, 128), (2, 256), (2, 512)),
        lr, num_classes)

ResNet18().layer_summary((1, 1, 96, 96))

```

```

Sequential output shape:      torch.Size([1, 64, 24, 24])
Sequential output shape:      torch.Size([1, 64, 24, 24])
Sequential output shape:      torch.Size([1, 128, 12, 12])
Sequential output shape:      torch.Size([1, 256, 6, 6])
Sequential output shape:      torch.Size([1, 512, 3, 3])
Sequential output shape:      torch.Size([1, 10])

```

```

In [ ]: model = ResNet18(lr=0.01)
trainer = d2l.Trainer(max_epochs=10, num_gpus=1)
data = d2l.FashionMNIST(batch_size=128, resize=(96, 96))
model.apply_init([next(iter(data.get_dataloader(True)))[0]], d2l.init_cnn)
trainer.fit(model, data)

```

```

In [ ]: class ResNeXtBlock(nn.Module): #@save
    """The ResNeXt block."""
    def __init__(self, num_channels, groups, bot_mul, use_1x1conv=False,
        strides=1):
        super().__init__()
        bot_channels = int(round(num_channels * bot_mul))
        self.conv1 = nn.Conv2d(bot_channels, kernel_size=1, stride=1)
        self.conv2 = nn.Conv2d(bot_channels, kernel_size=3,
            stride=strides, padding=1,
            groups=bot_channels//groups)
        self.conv3 = nn.Conv2d(num_channels, kernel_size=1, stride=1)
        self.bn1 = nn.BatchNorm2d()
        self.bn2 = nn.BatchNorm2d()
        self.bn3 = nn.BatchNorm2d()
        if use_1x1conv:
            self.conv4 = nn.Conv2d(num_channels, kernel_size=1,
                stride=strides)

```

```
        self.bn4 = nn.LazyBatchNorm2d()
    else:
        self.conv4 = None

    def forward(self, X):
        Y = F.relu(self.bn1(self.conv1(X)))
        Y = F.relu(self.bn2(self.conv2(Y)))
        Y = self.bn3(self.conv3(Y))
        if self.conv4:
            X = self.bn4(self.conv4(X))
        return F.relu(Y + X)
```

```
In [ ]: blk = ResNextBlock(32, 16, 1)
        X = torch.randn(4, 32, 96, 96)
        blk(X).shape
```

```
In [ ]:
```